



NPTEL ONLINE CERTIFICATION COURSES

Course Name: Deep Learning

Faculty Name: Prof. P. K. Biswas

Department : E & ECE, IIT Kharagpur

Topic

Lecture 60: Generative Adversarial Network

CONCEPTS COVERED

Concepts Covered

- Generative Model
- Intuitions behind VAE
- Variational Inference
- Practical Realization of VAE
- Generative Adversarial Network
- Applications of GAN



Variational Autoencoder : Variational Inference

□ Our initial objective: minimize $\text{KL}(Q(z|x) || P(z|x))$

□ Which is same as maximizing

$$\sum_z Q(z|x) \log \frac{P(x,z)}{Q(z|x)}$$



Variational Lower Bound

➤ So, aim now is: *maximize*

$$L = \sum_z Q(z|x) \log \frac{P(x,z)}{Q(z|x)} = \sum_z Q(z|x) \log \frac{P(x|z)P(z)}{Q(z|x)}$$



Variational Autoencoder : Variational Inference

Maximize

$$\begin{aligned} L &= \sum_z Q(z|x) \log \frac{P(x,z)}{Q(z|x)} = \sum_z Q(z|x) \log \frac{P(x|z)P(z)}{Q(z|x)} \\ &= \sum Q(z|x) \log P(x|z) + \sum Q(z|x) \log \frac{P(z)}{Q(z|x)} \end{aligned}$$



Variational Autoencoder : Variational Inference

Maximize

$$\begin{aligned} L &= \sum_z Q(z|x) \log \frac{P(x,z)}{Q(z|x)} = \sum_z Q(z|x) \log \frac{P(x|z)P(z)}{Q(z|x)} \\ &= \underbrace{\sum Q(z|x) \log P(x|z)}_{E_{Q(z|x)} \log P(x|z)} + \underbrace{\sum Q(z|x) \log \frac{P(z)}{Q(z|x)}}_{-KL(Q(z|x) || P(z))} \end{aligned}$$



Variational Autoencoder : Variational Inference

Maximize

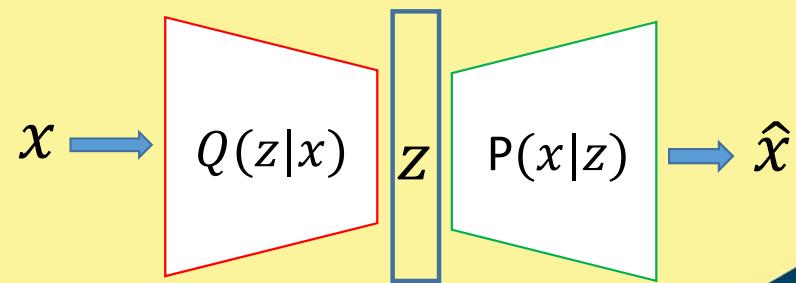
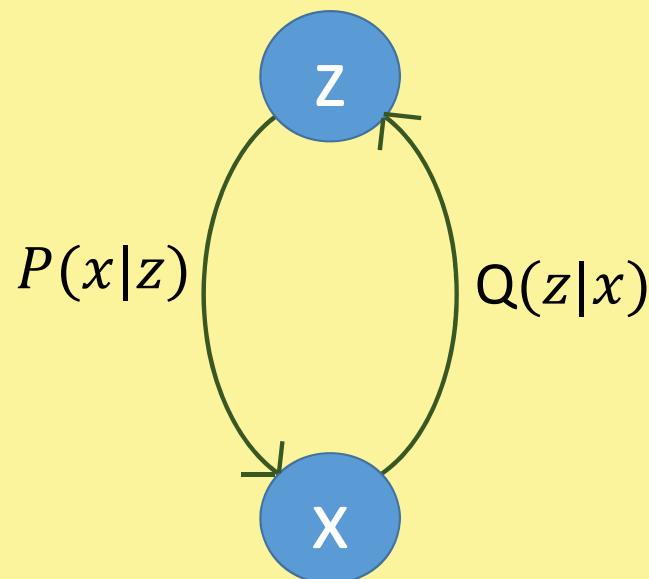
$$\begin{aligned} L &= \sum_z Q(z|x) \log \frac{P(x,z)}{Q(z|x)} = \sum_z Q(z|x) \log \frac{P(x|z)P(z)}{Q(z|x)} \\ &= \underbrace{\sum Q(z|x) \log P(x|z)}_{E_{Q(z|x)} \log P(x|z)} + \underbrace{\sum Q(z|x) \log \frac{P(z)}{Q(z|x)}}_{-KL(Q(z|x) || P(z))} \end{aligned}$$

- Translate the loss functions into an auto-encoder architecture.



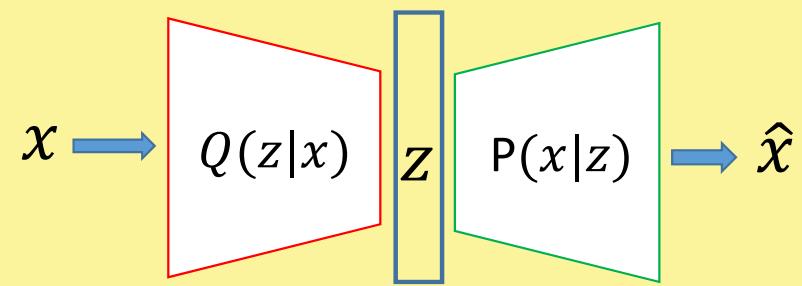
Variational Autoencoder : Network Realization

- We have the following graphical model
- Realize both $P(\cdot)$ and $Q(\cdot)$ with neural networks



Variational Autoencoder : Network Realization

- The z codes we get here should match with the distribution of $P(z)$ and we can decide what prior distribution to choose for $P(z)$.
- Usual practice is to select a Normal distribution $N(0, I)$ for the prior.

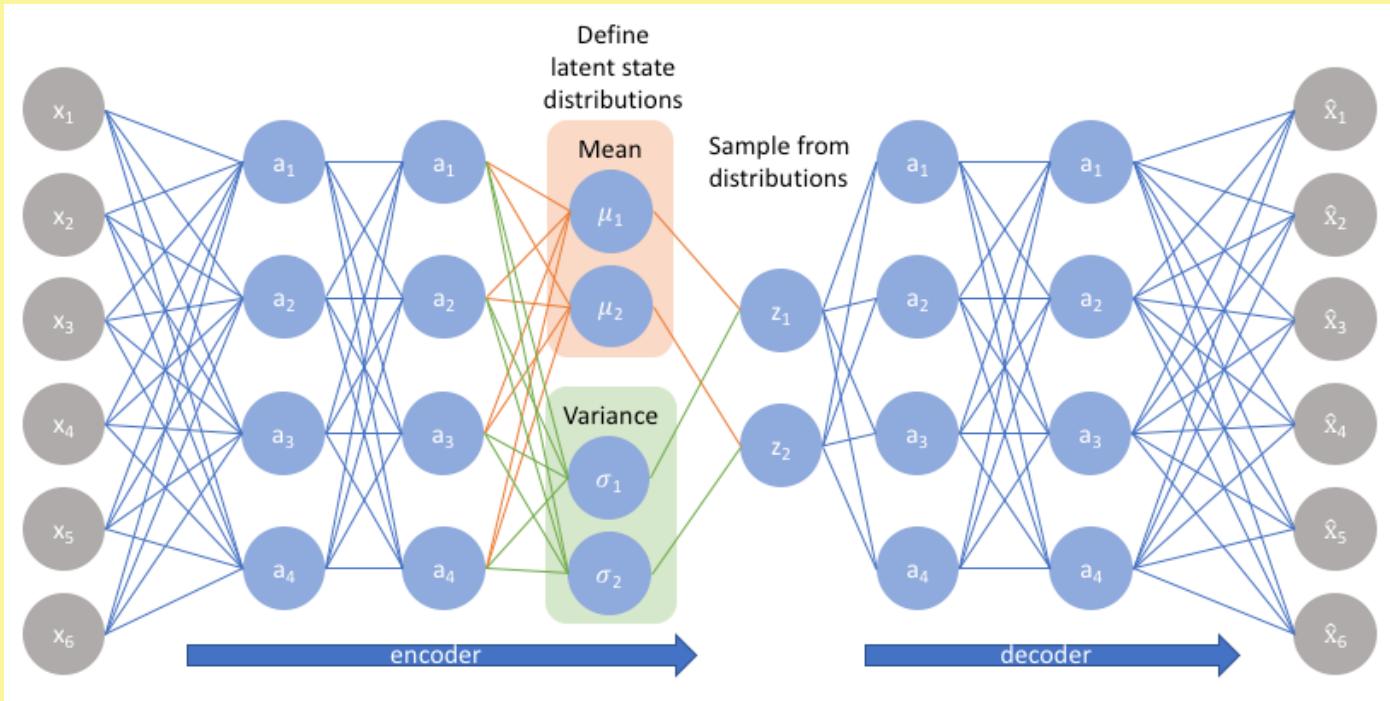


Variational Autoencoder : Network Realization

- Instead of generating a fixed code for an input, Encoder now gives parameters of the distribution of the latent code.
- For a given input x , we need to generate mean vector $\mu(x)$ and diagonal covariance matrix, $\Sigma(x)$.
- We need to SAMPLE a code from that latent distribution and pass forward to the Decoder.



Variational Autoencoder : Network Realization



<https://www.jeremyjordan.me/variational-autoencoders/>

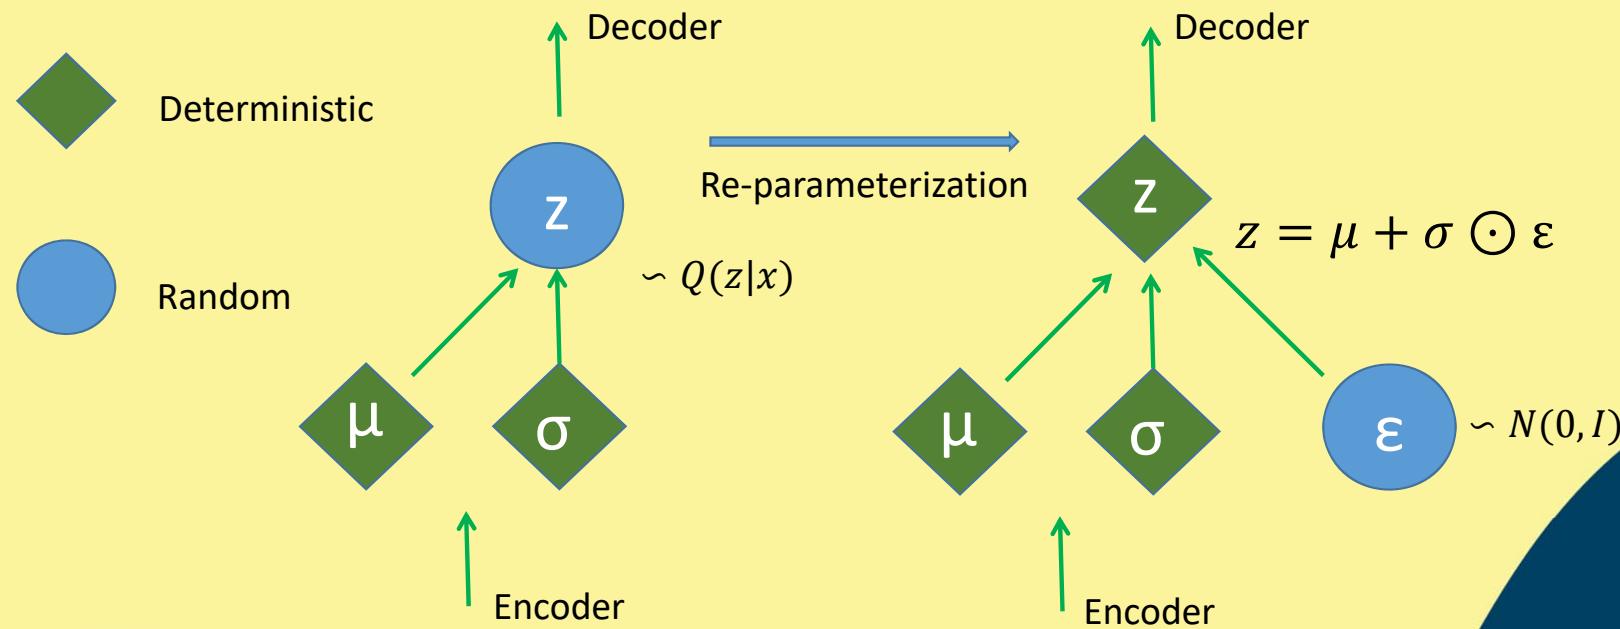
Variational Autoencoder : Network Realization

Sampling breaks computational graph
and
hinders Gradient Descent based
optimization

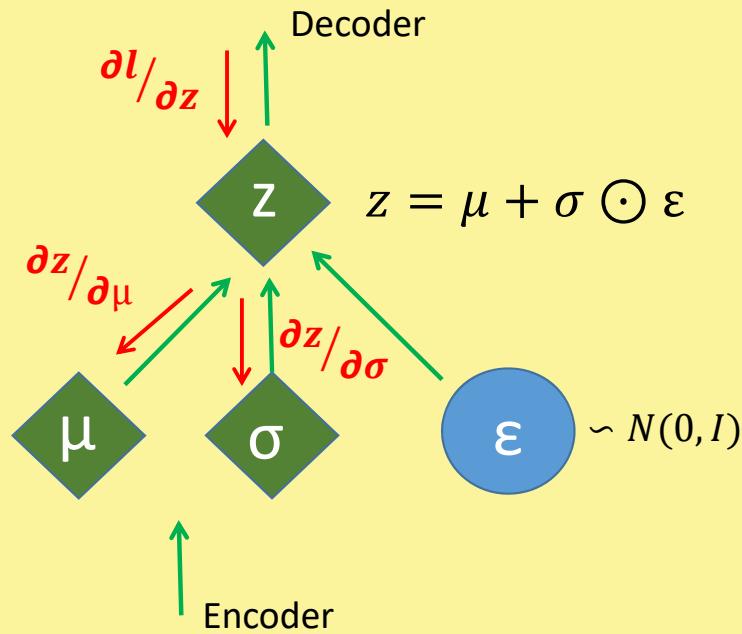


Variational Autoencoder : Reparameterization Trick

- We randomly sample ϵ from a unit Gaussian, and then shift the randomly sampled ϵ by the latent distribution's mean μ and scale it by the latent distribution's variance σ .



Variational Autoencoder : Reparameterization Trick



Re-parameterization enables

- Optimization of the parameters of the distribution.
- Still maintaining the ability to randomly sample from that distribution.



Variational Autoencoder : Coding the Cost Functions

$$E_{Q(z|x)} \log P(x|z) - KL(Q(z|x) || P(z))$$

Maximize

Minimize



Variational Autoencoder : Coding the Cost Functions

- Maximizing $E_{Q(z|x)} \log P(x|z)$ is a maximum likelihood estimation.
It is observed all the time in discriminative supervised model, for example Logistic Regression, SVM, or Linear Regression.
- In the other words, given an input z and an output x , we want to maximize the conditional distribution $P(x|z)$ under some model parameters.
- So we could implement it by using any classifier with input z and output x , then optimize the objective function by using for example log loss or regression loss.



Variational Autoencoder : Coding the Cost Functions

- We want to minimize the second component of the loss, $KL((Q(z|x) || P(z))$
- We assumed that $P(z)$ follows $N(0, I)$, so we have to push $Q(z|x)$ towards $N(0, I)$

Assuming $P(z)$ to be $N(0, I)$ has 2 advantages:

- Easy to sample latent vectors from $N(0, I)$ when we want to generate samples.
- Assuming $Q(z|x)$ to be a Gaussian distribution with parameters, $\mu(x)$ and $\Sigma(x)$ allows $KL(Q(z|x) || P(z))$ to be in a closed form and easy for optimization.



Variational Autoencoder : Coding the Cost Functions

$$KL(N(\mu(x), \Sigma(x)) || N(0, I)) = 0.5 * [tr(\Sigma(x)) + \mu(x)^T \mu(x) - k - \log \det(\Sigma(x))]$$

- k is dimension of the latent code
- $tr(\Sigma(x))$ is trace of a covariance matrix
- $\Sigma(x)$ is the diagonal covariance matrix. So, its determinant can be computed as product of its diagonal entries.
- In practice $\Sigma(x)$ can be predicted only as vector containing the diagonal entries



Variational Autoencoder : Coding the Cost Functions

$$KL(N(\mu(x), \Sigma(x)) || N(0, I) = 0.5 * [tr(\Sigma(x) + \mu(x)^T \mu(x) - k - \log \det(\Sigma(x)))]$$

$$= 0.5 * \left[\sum_k \Sigma(x)_k + \sum_k (\mu(x)_k)^2 + \sum_k 1 - \log \prod_k \Sigma(x)_k \right]$$

$$= 0.5 * \left[\sum_k \Sigma(x)_k + \sum_k (\mu(x)_k)^2 + \sum_k 1 - \sum_k \log \Sigma(x)_k \right]$$

$$= 0.5 * \sum_k [\Sigma(x)_k + (\mu(x)_k)^2 + 1 + \log \Sigma(x)_k]$$



Variational Autoencoder : Coding the Cost Functions

In practice, we predict $\log \Sigma(x)$ instead of only $\Sigma(x)$ since it is numerically better to exponentiate a value during run time rather than taking log.

$$KL(N(\mu(x), \Sigma(x)) \parallel N(0, I))$$

$$= 0.5 * \sum_k [\exp(\Sigma(x)_k) + (\mu(x)_k)^2 + 1 + \log \Sigma(x)_k]$$



Variational Autoencoder :After training

Visualizing Reconstructions:



Input

Reconstructions



Variational Autoencoder :After training

Visualizing Reconstructions:



Input

Reconstructions

Using as Generative Model

- Sample a random vector from $N(0, I)$
- Feed forward the vector through the pre-trained Decoder



Generated Samples



Variational Autoencoder : Generative Model



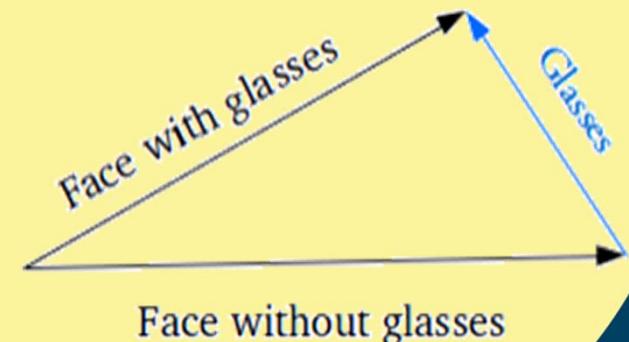
VAE generating novel faces after trained
on CelebA dataset



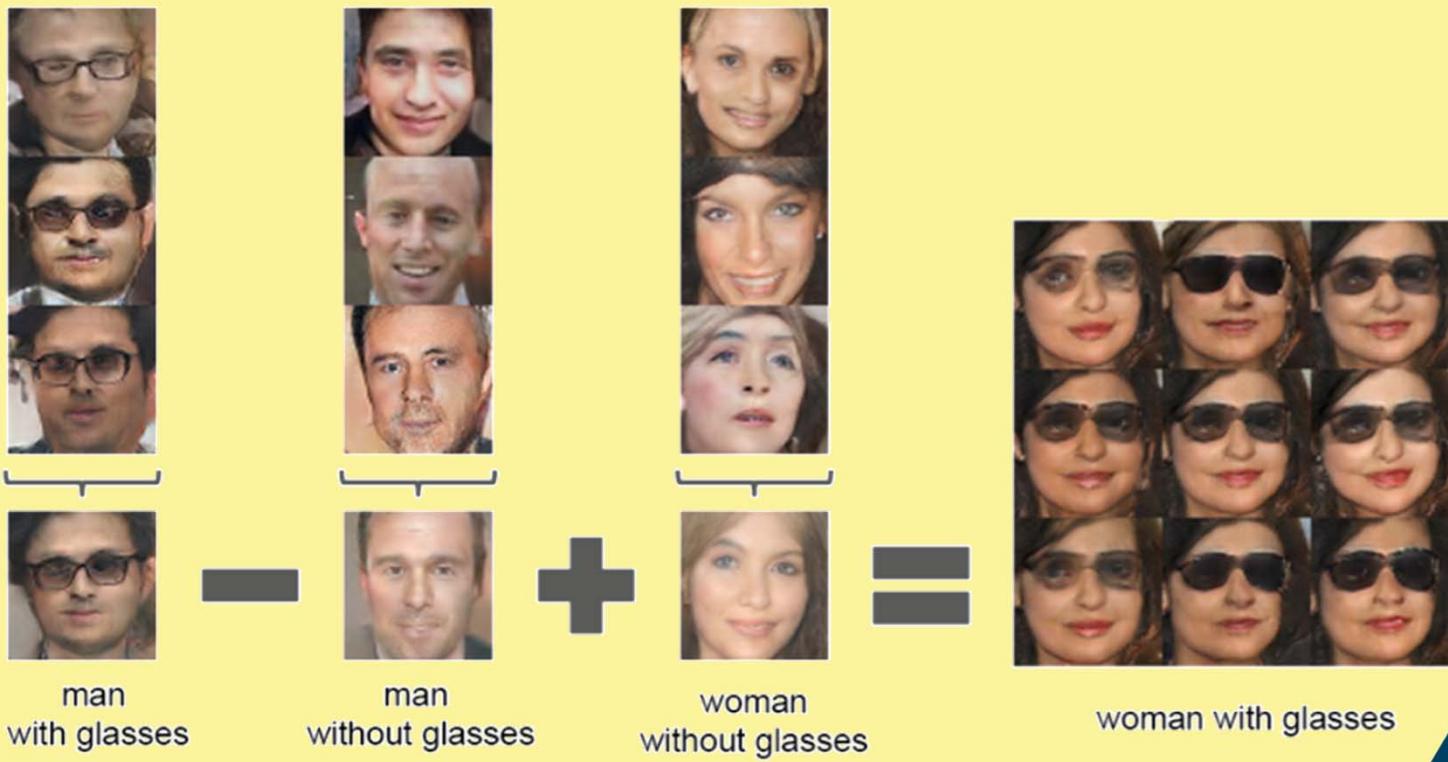
Variational Autoencoder : Vector Arithmetic

How do you interpolate between two samples ?

- Take a face image with glasses and find the latent code (C_1)
- Take another face without glasses and find latent code (C_2)
- $C_3 = C_1 - C_2$ gives code for glasses
- Take a new face without glasses and find latent code (C_4)
- $C_3 + C_4$ will overlay glasses on this new image
- Such transitions are possible only if the latent space is continuous instead of clusters



Variational Autoencoder : Generative Model



Generative Adversarial Network (GAN)

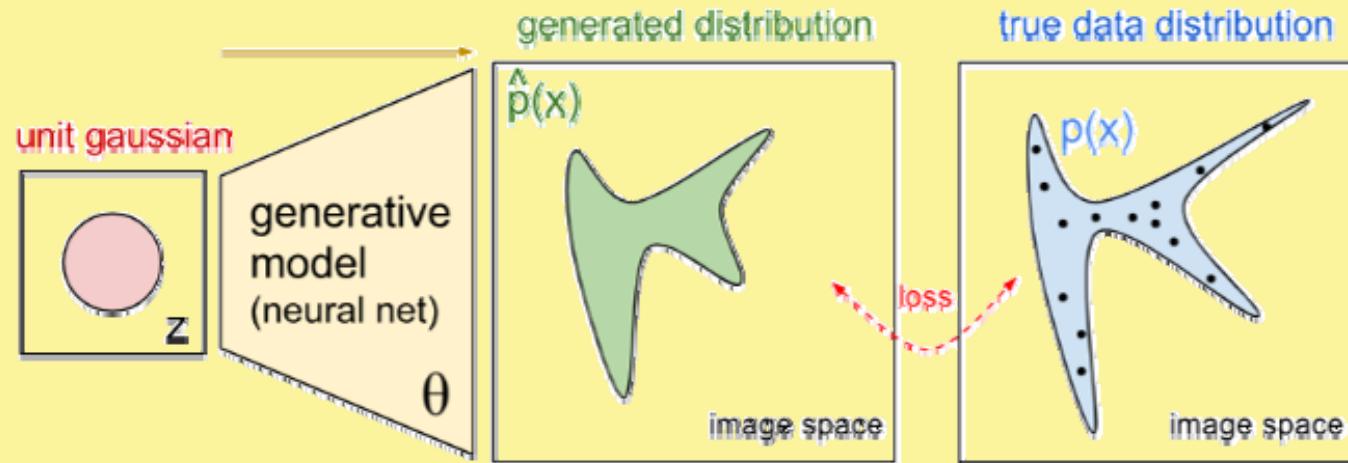


Implicit Generative Models

- Implicitly defines a probability distribution.
- Sample code vector, z , from a simple and fixed distribution (e.g. spherical Gaussian or Uniform).
- A generator network is trained as a differentiable network to map z to a data point x .



Implicit Generative Models



<https://openai.com/blog/generative-models/>

Implicit Generative Models

- ❑ Blue Region shows areas with high probability of real image.
- ❑ Black dots represent actual images from true distribution $p(x)$.
- ❑ Generative model (parameterized by θ) also describes a function $\hat{p}(x)$
 - Takes points (latent codes) from an unit Gaussian distribution.
 - Maps those points to a generator distribution.
 - θ can be optimized to reduce $KL(p(x)||\hat{p}(x))$
 - Green distribution starts randomly then aligns with blue distribution



<https://openai.com/blog/generative-models/>

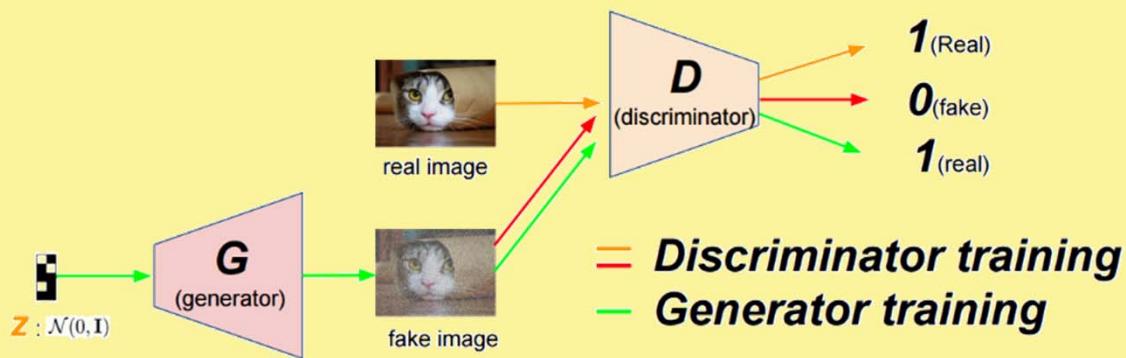
GAN Overview

- In GAN the main idea is to have two neural networks compete with each other.
- Its Game Theoretic Approach.
 - **Generator** network samples a z vector and tries to produce realistic samples.
 - **Discriminator** network tries to distinguish fake samples (from Generator) and real samples.

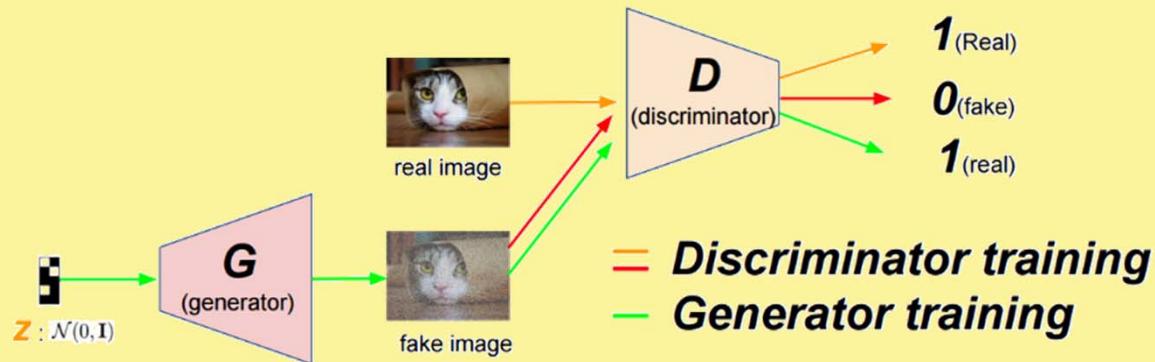


GAN Overview

- Assume $D(x)$ represents probability of belonging to real class for a given sample, x
- Discriminator will try to increase $D(x)$ for real samples and decrease $D(x)$ for fake/generated samples
- Generator will try to increase $D(x)$ for generated samples



GAN Overview

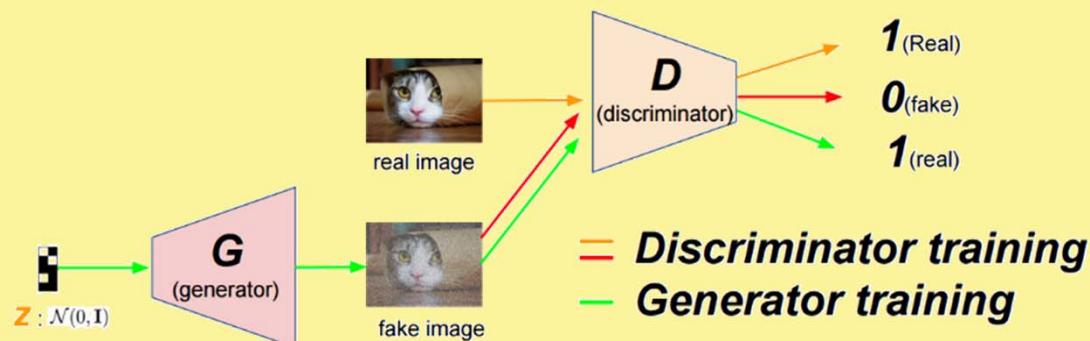


Training the Discriminator

$$\max_D V(D, G) = \underbrace{E_{x \sim p_{data}(x)} \log D(x)}_{\text{Maximize probability for real}} + \underbrace{E_{z \sim p_z(z)} [\log\{1 - D(G(z))\}]}_{\text{Minimize probability for generated}}$$



GAN Overview



Training the Generator

$$\begin{aligned} \min_G V(D, G) &= E_{z \sim p_z(z)} [\log\{1 - D(G(z))\}] \\ &\equiv \max_G E_{z \sim p_z(z)} [\log D(G(z))] \end{aligned}$$

Maximize probability for generated



GAN Training : Alternate updates of D and G

for number of training iterations **do**

for k steps **do**

- Sample minibatch of m noise samples $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$ from noise prior $p_g(\mathbf{z})$.
- Sample minibatch of m examples $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$ from data generating distribution $p_{\text{data}}(\mathbf{x})$.
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(\mathbf{x}^{(i)}) + \log (1 - D(G(\mathbf{z}^{(i)}))) \right].$$

end for

- Sample minibatch of m noise samples $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$ from noise prior $p_g(\mathbf{z})$.
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(\mathbf{z}^{(i)}))).$$

end for



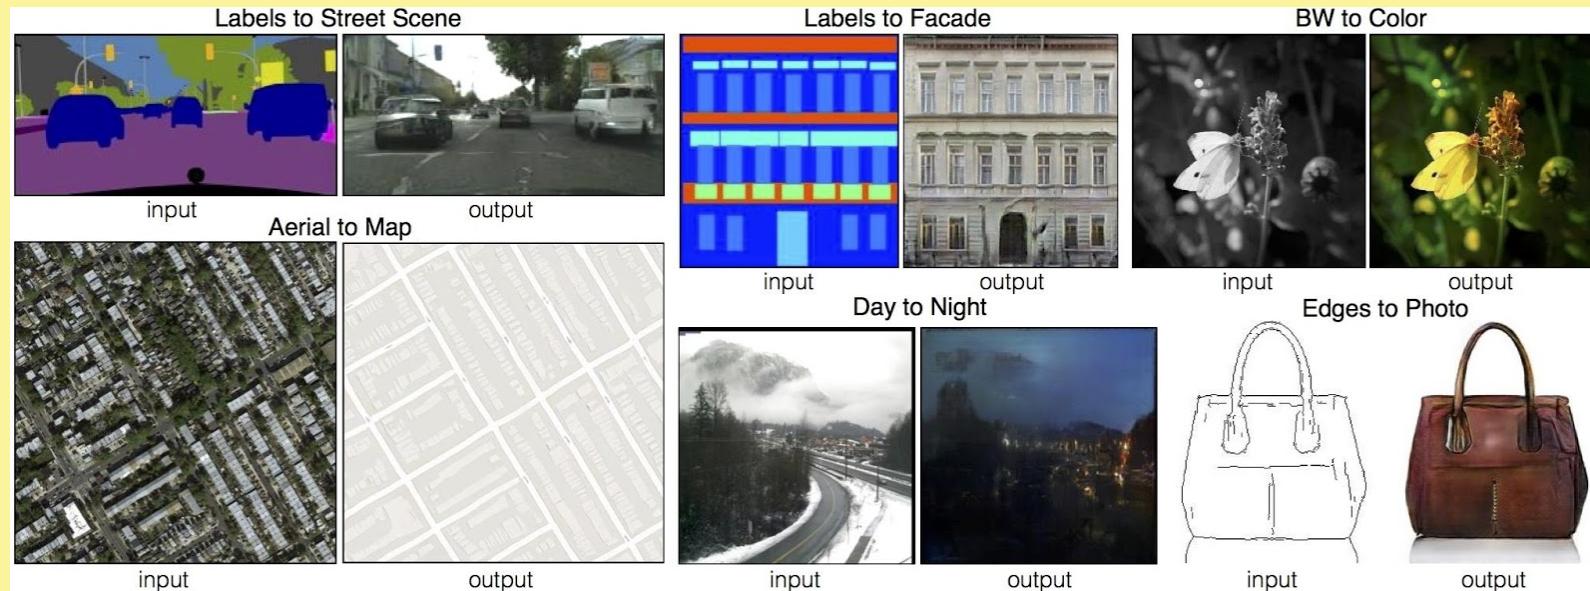
Goodfellow et al. "Generative Adversarial Networks",
NeurIPS, 2014

GAN Applications: High Resolution Image Synthesis



Karras, Tero, Timo Aila, Samuli Laine, and Jaakko Lehtinen.
"Progressive growing of gans for improved quality, stability, and variation." ICLR, 2018.

GAN Applications: Image to Image Translation



*Isola, Phillip, Jun-Yan Zhu, Tinghui Zhou, and Alexei A. Efros.
"Image-to-image translation with conditional adversarial networks." *CVPR*, 2017

GAN Applications: Video to Video Translations



Input Labels



Style 1



Style 2



Style 3



Wang, Ting-Chun, Ming-Yu Liu, Jun-Yan Zhu, Guilin Liu, Andrew Tao, Jan Kautz, and Bryan Catanzaro. "Video-to-video synthesis." *NeurIPS*, 2018

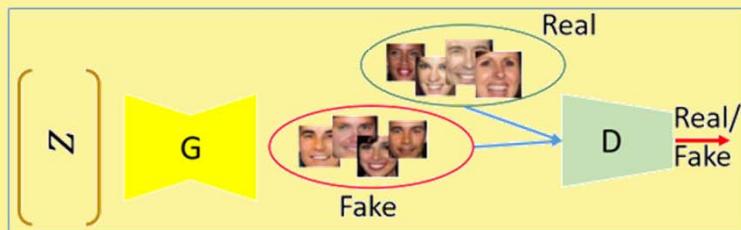
GAN Applications: Image Inpainting

Input: Masked/damaged image, I_d , with binary mask, M

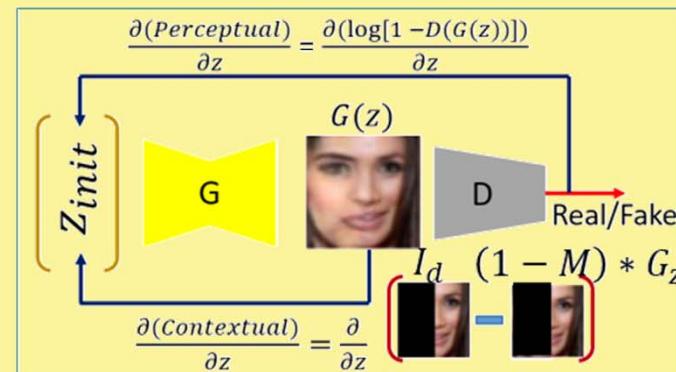
Intermediate Output : Image, I_G , after iterative optimization for z

Final Output: Inpainted image, $\hat{I}_d = M * I_d + (1 - M) * I_G$

Stage 1: Pre-training a GAN



Stage 2: Iterative search for z



Yeh, R. A., Chen, C., Yian Lim, T., Schwing, A. G., Hasegawa-Johnson, M., & Do, M. N. (2017). Semantic image inpainting with deep generative models. In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR).

GAN Applications: Image Inpainting

PROPOSED IMAGE INPAINTING (5X SPEEDUP)

Nearest Neighbour Search for Better z Initialization

- Sample a pool of z vectors and pass through pre-trained G
- Data + Structure loss = $L_{nn}(\cdot)$ between masked, I_d & pooled, I_p^i
- Select z_{init} as initial solution, s.t: $z_{init} = \underset{z^i}{\operatorname{argmin}} L_{nn}(I_d, G(z^i))$

Data Loss, L_D

$$L_D^i = |I_d - M * p_i|$$

Structure Loss, L_S

$$L_S^i = |\Delta_x I_d - \Delta_x M * I_p^i| + |\Delta_y I_d - \Delta_y M * I_p^i|$$

Both uses only
unmasked pixels info



Lahiri, A., Jain, A. K., Nadendla, D., & Biswas, P. K., "Faster Unsupervised Semantic Inpainting: A GAN Based Approach", ICIP 2019.

GAN Applications: Image Inpainting



Lahiri, A., Jain, A. K., Nadendla, D., & Biswas, P. K., "Faster Unsupervised Semantic Inpainting: A GAN Based Approach", ICIP 2019.

GAN Applications: Video Inpainting

- PROPOSED VIDEO INPAINTING (80X SPEEDUP)

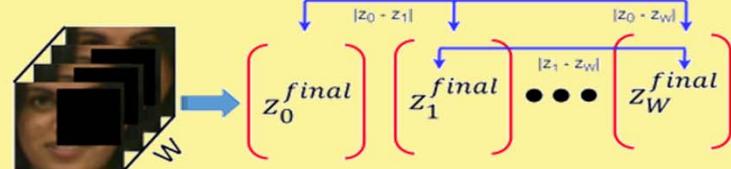
• Reuse Noise Priors

- Exploit temporal redundancy
- Temporal neighbours should have close z representaitons



• Group Consistency Loss

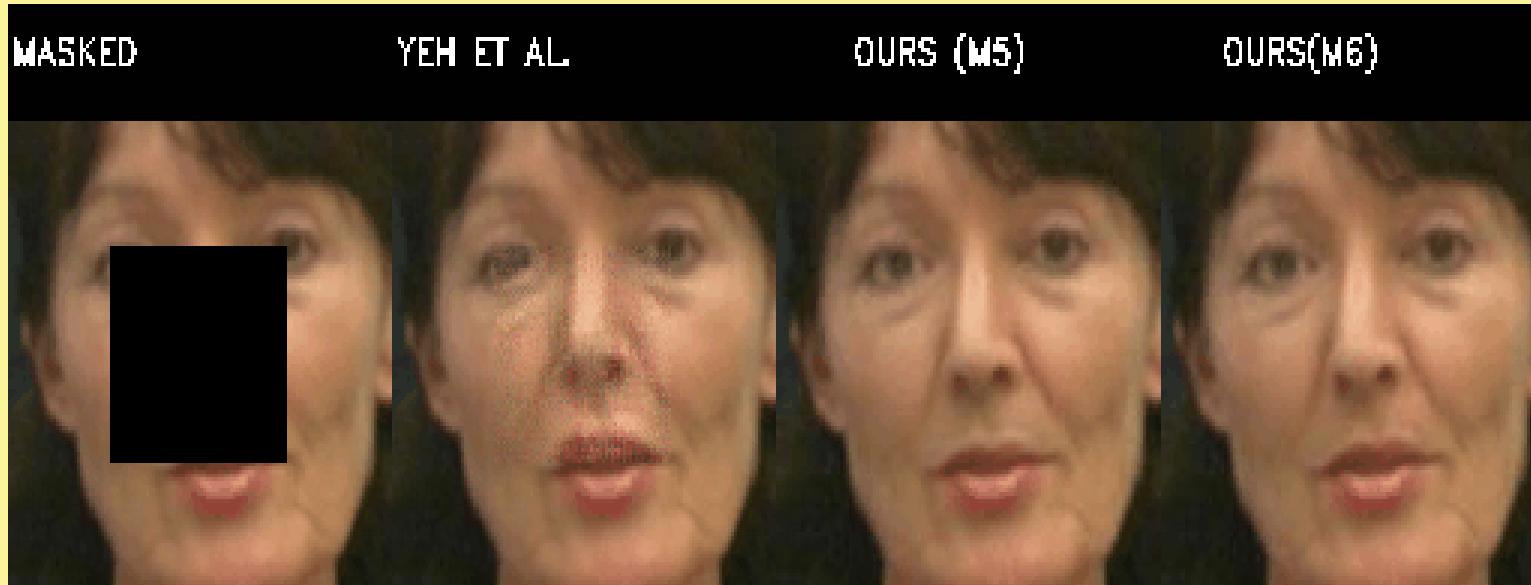
- Penalize if a local temporal neighbourhood of W frames differ
- Helps in reducing sudden flickering effects across frames
- $Loss = |z_k - z_j| \forall i \in [1, 2, \dots, W]; \forall j \in [1, 2, \dots, W]$



Lahiri, A., Jain, A. K., Nadendla, D., & Biswas, P. K., "Faster Unsupervised Semantic Inpainting: A GAN Based Approach", ICIP 2019.



GAN Applications: Video Inpainting



Lahiri, A., Jain, A. K., Nadendla, D., & Biswas, P. K., "Faster Unsupervised Semantic Inpainting: A GAN Based Approach", ICIP 2019.



NPTEL ONLINE CERTIFICATION COURSES

*Thank
you*